

Part A – Product Review

Overall Assessment

The PRD clearly explains the value of Custom Reporting and the main user flows. However, before release, I would want to close several gaps that could affect data access, report accuracy, reliability, and the user's ability to trust reports over time.

I prioritized the following five gaps based on their potential impact on customers and on whether they could block confident implementation and testing.

1. Report Visibility, RBAC, and Partial Data Access Are Not Fully Defined

What is missing or unclear

The PRD allows reports to be set as “Only for me” or “Organization level”, and states that report CRUD operations are controlled by the platform's RBAC model. However, it does not clearly define how access to the report itself interacts with access to the underlying data.

For example:

- What happens when a user with more limited permissions opens or downloads an organization-level report created by someone with broader access?
- How should the report behave when different viewers or recipients are authorized to see only subsets of the underlying data, such as specific security modules, sites, regions, or environments?
- Should the system generate a different version of the report for each user, hide restricted sections, display those sections as unavailable, show a permission-related error, or block access to the report entirely?
- What actions are organization-level users allowed to perform: view, download, edit, delete, or connect the report to a Workflow?

Why it matters

Reports may contain sensitive security information. A user should never gain access to data through a report that they would not normally be allowed to access in the platform.

Partial permissions also create a product and technical challenge. If different recipients are allowed to see different parts of the same report, the system may need to generate personalized versions rather than a single shared file. Without a clear model, this could create security and compliance risks as well as inconsistent user experiences.

What needs to be decided

The team should define:

- Whether permissions are evaluated based on the report creator, the current viewer or recipient, or both.
- How the report behaves when a viewer or recipient is authorized to see only part of the underlying data.
- Whether this results in per-user report generation, partial/redacted reports, a permission-related error, or full access denial.
- Which actions organization-level users are allowed to perform.
- Whether authorization is checked again when a report is generated or downloaded.
- What happens when a user's permissions change after a report has already been created.

2. The Relationship Between a Saved Report and Its Source Is Unclear

What is missing or unclear

Reports can depend on existing objects in the platform, such as Dashboards and Saved Views. For Custom Dashboards, the PRD introduces "Sync board changes to future reports", but it does not fully define what the report stores or how dependent it remains on its original source.

For example:

- What happens if the source Dashboard is deleted?
- What happens if a Saved View used as the report scope is edited or deleted?
- If Sync is disabled, does the report keep a snapshot of the original Dashboard configuration?
- What happens if a widget, filter, or field used by the report no longer exists?
- What happens if the report owner later loses access to the source?

Why it matters

Reports are intended to be reusable objects that customers can depend on over time. If a report unexpectedly changes or stops working because another object in the system changed, users may receive incorrect or incomplete information without realizing it.

What needs to be decided

- What information is copied into the report versus kept linked to the source.
- How source modifications should affect existing reports.
- What should happen if the source becomes inaccessible.
- How a report should behave if one of its referenced elements, such as a widget, filter, or field, is no longer available.
- What should happen when Sync is disabled and the source later changes.

3. Scheduled Delivery Needs a Clearer Authorization, Dependency, and Failure Model

What is missing or unclear

The PRD allows a saved report to be connected to the Workflow engine and delivered automatically by email. Large files are delivered through an authenticated download link.

However, several important parts of this flow are not defined.

For example:

- Can reports be sent to any email address or only to users within the organization?
- Are external recipients or distribution lists supported?
- Who is allowed to access an authenticated download link, and how long is it valid?
- What happens if report generation succeeds but email delivery fails?
- Who is notified when a scheduled delivery fails?
- What happens if a report connected to an active Workflow is edited or deleted?
- Can multiple Workflows use the same report?

Why it matters

Automation is one of the main benefits of Custom Reporting. Customers may rely on reports being delivered without manually checking each execution.

A silent failure could mean important information never reaches its intended audience. At the same time, unclear recipient and authorization rules could expose sensitive information outside the intended audience.

There is also a dependency between the Report object and the Workflow that needs to remain valid over time.

What needs to be decided

The team should define:

- Supported recipient types and authorization rules.
- Authentication and expiration rules for large-report download links.
- Retry and notification behavior when generation or delivery fails.
- Whether users can view delivery status or history.
- What happens to an active Workflow when its report is edited, deleted, or becomes invalid.

4. Behavior Beyond the 500K Record Limit Is Not Defined

What is missing or unclear

The CSV creation mockup states that the CSV format supports up to 500K records, which provides an important system limit. However, the PRD does not define what happens when the selected scope returns more than 500K records.

In addition, Summary Reports support a relative time range, while the Operational Data flow does not define a similar time-range option. This is particularly relevant for CSV reports, where limiting the reporting period could help control the amount of data included in an export that is already limited to 500K records.

For example:

- Is report generation blocked?
- Are only the first 500K records exported?
- If the data is truncated, how are those records selected and ordered?
- Can the export be split into multiple files?
- Is the user asked to narrow the report scope?
- Does the user receive a warning before starting generation?
- Can users limit a CSV report by time range?
- What happens if a scheduled report that previously contained fewer than 500K records grows beyond the limit over time?
- Does the scheduled delivery fail, truncate the data, or notify the report owner?

Why it matters

Enterprise environments can contain large volumes of findings, so exceeding the 500K limit is a realistic scenario.

If the system silently truncates data, users may assume the report is complete when it is not. This therefore affects both data accuracy and reliability.

What needs to be decided

The team should define:

- The exact behavior when a report exceeds 500K records.
- Whether truncation is ever allowed and, if so, how it is clearly communicated.
- Whether users should be required to narrow their scope or whether multiple-file exports are supported.
- Whether CSV reports should support a time range
- How scheduled reports behave when they cross the limit.

5. The Written PRD and Figma Mockups Are Not Fully Aligned

***If the mockups are mainly intended to illustrate the general direction rather than define the final behavior, this point becomes less significant. In that case, it would still be helpful to clarify which source should be treated as the final reference for implementation and testing.**

What is missing or unclear

There are several differences between the written product requirements and the Figma mockups, making it unclear which definition represents the intended release.

For example:

- The written PRD lists eight supported security modules for Operational Data reports, while the CSV creation mockup presents only five data-table options.
- The written PRD refers to the PDF report type as “Summary Report,” while the mockup uses “Executive Report.” The PRD itself also notes that the final naming has not yet been decided.
- The written creation flow contains four main steps, while the mockups introduce additional report-type-specific configuration steps.
- The PRD states that an existing Saved View can be selected as the report scope, but this option is not shown in the provided creation mockups.
- The written requirements state that users have control over both column selection and ordering for Operational Data reports, while the mockup does not clearly show how column ordering is configured.

Why it matters

Engineering, QA, and Design need a single source of truth for what is included in the release.

If the written requirements and the mockups describe different behavior, different teams may implement and test against different assumptions. This can lead to missing functionality, inconsistent UX, incorrect test coverage, or late changes before release.

What needs to be decided

The team should define:

- The final list of supported modules and data tables for Operational Data reports.
- The final names of the report types.
- Whether Saved Views are supported as a scope in the initial release.
- Whether column ordering is part of the initial release and how it is configured.
- Which report-creation steps are required for each report type.

Additional Lower-Priority Questions

I would not prioritize these above the five issues above, but I would still clarify them during implementation and testing:

- What happens when the selected scope or filters return no data? Is an empty report generated, is generation blocked, or is it handled differently for scheduled reports?
- Which fields are mandatory for each report type, and what happens if a user attempts to create a report without providing a required input?
- How should an Operational Data Report behave if one of its selected fields is later removed or renamed?
- Are duplicate report names allowed, and if so, how are reports distinguished when referenced from other areas such as Workflows?
- Which report actions are included in the audit log beyond create, update, and delete?
- What happens to personal reports when their creator leaves the organization?
- Is Operational Data limited to a single security module across all report creation entry points, including "Save as Custom Report"?
- How exactly are relative time ranges calculated, including timezone and scheduled execution behavior?

Part B - Test Plan

Test Approach

The goal of this test plan is to check the main Custom Reporting flows, make sure each report matches what the user selected, verify permissions, and cover the main edge cases from Part A.

When the PRD does not define the exact behavior, I do not assume a solution. I define what must be true and mark the test as Partial or Blocked until the open decision is closed.

Status and Priority

- Ready: the PRD is clear enough to run the test and decide pass or fail.
- Partial: the main outcome can be tested, but one detail still needs a decision.
- Blocked: a final pass or fail result depends on an open decision from Part A.
- Critical: must pass before release because of security, tenant isolation, or silent data risk.
- High: a major issue in a core flow or in reliability.
- Medium: an important edge case with lower impact or a reasonable workaround.
- Low: a usability or presentation issue that does not block release.

Open Decisions

- Gap 1: RBAC rules for report data, partial access, scheduled runs, and recipient permissions.
- Gap 2: what a saved report keeps from its Dashboard or Saved View, and what happens when the source changes.
- Gap 3: recipient rules, download link access and expiry, delivery failures, Workflow dependency, and file retention.
- Gap 4: what happens above 500K records and what time period a CSV represents.
- Gap 5: final alignment between the PRD and Figma, including modules, report names, Saved Views, column ordering, and creation steps.
- Other open points: exact time range rules, zero result behavior, report naming, and audit log scope.

#	Scenario	Expected Result	Priority	Status
1. Report Creation and Data Integrity				
1	Summary Report from a Predefined Dashboard Create a Summary Report with valid details, scope, and time range.	A readable PDF is created from the selected Dashboard and includes only the selected scope and time period.	High	Ready
2	Summary Report from a Custom Dashboard Create a Summary Report from an existing Custom Dashboard.	A readable PDF is created from the selected Custom Dashboard and reflects the selected scope and time period.	High	Ready
3	Create an Operational Data Report Create a CSV using a supported module, scope, and selected columns.	The CSV includes only matching records and the selected columns. If column ordering is part of this release , it must follow the final requirement.	High	Partial (Gap 5: column ordering)
4	Report Matches the Source Data Apply a known filter or scope, generate a report, and compare the rows, values, and key totals with the source screen or Dashboard.	The report matches the same result set the user selected. Rows, filters, or totals must not change without explanation.	Critical	Ready
5	Generate a Saved Report Again Save a report, return to the Reports page later, and generate it again.	The saved configuration is kept and the same report can be generated again without rebuilding it.	Medium	Ready
6	Use a Saved View as Report Scope Create a report using an existing Saved View as the scope.	If Saved Views are included in this release, the report uses the same scope and filters as the selected Saved View.	Medium	Blocked (Gap 5: Saved View support)
7	Time Range Boundaries Generate Summary Reports around day or month boundaries, scheduled run times, and relevant timezones.	The report uses one clearly defined time window. The exact timezone and boundary rules must be defined before release.	Medium	Blocked (Time range rules)
2. Save as Custom Report				

#	Scenario	Expected Result	Priority	Status
8	Create a Report with Save as Custom Report Start from a supported table or board and choose Save as Custom Report.	The creation flow opens with the relevant settings already filled in, so the user does not need to rebuild the source filters.	High	Ready
9	Save as Custom Report Matches the Source View Create a report from a filtered table or view and compare the output with what was visible at the source.	The generated report matches the source view the user chose to save. Pre-filled settings are not enough if the final output is different.	Critical	Ready
10	Module Selection Is Consistent Compare Operational Data creation from the regular flow and from Save as Custom Report, including a source that may include more than one module.	The same module rule applies in both entry points. The exact rule for one or multiple modules must be defined first.	Medium	Blocked (Gap 5: module rule)
11	User Without Create Permission Use a user without report create permission and try both Create Report and Save as Custom Report.	Both attempts are blocked. Save as Custom Report must not bypass report creation permissions.	High	Ready
3. Report Management and UX				
12	Edit a Saved Report Edit an existing saved report and generate it again.	The changes are saved to the intended report and are used the next time it is generated.	Medium	Ready
13	Reports Page Search, Filters, and Empty State Use the supported search and filters, and also open the Reports page when there are no saved reports.	Search and filters show the correct reports. The empty state is clear and does not look like an error or endless loading.	Low	Ready

#	Scenario	Expected Result	Priority	Status
14	Duplicate Report Names Create two reports with the same name, then search, download, edit, or connect one to a Workflow.	The system must not act on the wrong report because the names match. Whether duplicate names are allowed is a product decision.	Medium	Partial (Naming rules)
15	Long or Special Report Names Use a long report name and non Latin characters. Check the Reports page, file name, and scheduled email.	The name does not break the UI, file name, or email, and it does not make the report ambiguous. Exact limits can follow the final UX decision.	Low	Partial (Naming UX)
16	Required Inputs and Cancel Try to continue without required values such as report name or Dashboard or module. Also cancel before saving.	A report with missing required information cannot be saved as valid. Cancelling does not leave an unintended saved report.	Medium	Ready
17	Long Generation and Error State Trigger a slow generation and a failed generation. Check loading, repeated clicks, and the final state.	A running or failed job is not shown as successful, and repeated clicks do not create accidental duplicate generations. Exact messages can follow the final UX.	Medium	Partial (UX and failure handling)
4. Visibility, RBAC, and Tenant Isolation				
18	Only for Me Report Access Create a personal report and try to access it from another user, through the UI and by direct report or download reference where possible.	The second user cannot view, generate, or download the report or its data. Access control must not rely only on hiding the UI.	Critical	Ready
19	Organization Report with Partial Data Access Create an organization level report as a user with broad access, then open it as a user who can see only part of the underlying data.	The second user never sees data outside their own permissions. The exact handling, such as partial content, an error, or denial, follows the final RBAC model.	Critical	Blocked (Gap 1: partial data access)

#	Scenario	Expected Result	Priority	Status
20	Organization Report Action Permissions Try view, download, edit, delete, and connect to Workflow with users who have different roles. Use the UI and direct object access where possible.	Each action follows the final permission rules. Hiding a button is not enough if the action is still available another way.	High	Blocked (Gap 1: action permissions)
21	Permissions Change After Report Creation Create a report, then reduce or remove the creator or viewer permissions before a later generation or download.	A saved report must not keep access that the user no longer has. The exact permission check follows the final RBAC model.	High	Blocked (Gap 1: permission recheck)
22	Permissions for a Scheduled Report Run a scheduled report when no user is logged in.	The scheduled run uses one defined identity and never includes data outside that identity's permissions. The identity must be decided before release.	High	Blocked (Gap 1: scheduled identity)
23	Recipients with Different Data Permissions Send the same scheduled report to recipients who are allowed to see different parts of the data.	Each recipient sees only data they are allowed to access. The exact solution is not assumed and must follow the final permission model.	Critical	Blocked (Gap 1: recipient permissions)
24	Cross Tenant Access From Tenant B, try to list, open, download, or directly reference a report or generated file from Tenant A.	Tenant B cannot access report details or data from Tenant A through the UI or by direct reference.	Critical	Ready
25	Report Creator Is Deactivated Deactivate the creator of a personal report and of a scheduled report. Then try access and the next scheduled run.	Deactivation must not silently keep old permissions or continue delivery under an undefined owner. The final lifecycle rule must be defined.	Medium	Blocked (Gap 1 and Gap 3)
5. Source Changes and Sync				
26	Sync Enabled Create a report with Sync enabled, change the source Dashboard, and generate it again.	The next report reflects the updated Dashboard, as described in the PRD.	Medium	Ready

#	Scenario	Expected Result	Priority	Status
27	Sync Disabled Create a report with Sync disabled, change the source Dashboard, and generate it again.	The report follows one defined rule and does not mix old and new Dashboard settings. The exact rule must be decided before release.	Medium	Blocked (Gap 2: saved source behavior)
28	Source Dashboard Is Unavailable Delete the source Dashboard or remove the report owner's access, then generate the saved report.	The system does not silently create a misleading report from an unavailable source. The exact behavior follows the final source lifecycle rule.	High	Blocked (Gap 2: source lifecycle)
29	Saved View Is Edited or Deleted Use a Saved View as scope, then edit it and later delete it before generating again.	The report follows one defined relationship to the Saved View and does not silently use the wrong scope.	Medium	Blocked (Gap 2 and Gap 5)
30	Referenced Widget, Filter, or Field Is Missing Remove or rename a widget, filter, or field used by a saved report, then generate the report again.	The report does not silently omit or replace the missing item while still looking complete. Exact handling follows the final source rule.	High	Blocked (Gap 2: referenced items)
6. Scheduled Delivery and Generated Files				
31	Scheduled Email Delivery Connect a saved report to a Workflow and trigger a scheduled run. Check the email and attachment or link.	The correct report reaches the correct recipients at the scheduled time. The email identifies the right report and does not reveal extra data.	High	Partial (Email content not defined)
32	External Recipient or Distribution List Try scheduled delivery to an external email address and to a distribution list.	The final recipient policy is enforced. Unsupported or unauthorized recipients do not receive report data.	Medium	Blocked (Gap 3: recipient policy)
33	Large Report Download Link Trigger the large file flow. Try the link as the intended recipient, an unauthorized user, a user from another tenant, and after expiry.	Unauthorized, cross tenant, or expired access is rejected. The exact authentication and expiry rules must be defined before release.	Critical	Blocked (Gap 3: link access and expiry)

#	Scenario	Expected Result	Priority	Status
34	Generation or Email Delivery Failure Test a generation failure and a case where generation succeeds but email delivery fails. Check any status or history shown to the user.	A failed generation or delivery is never shown as successful. The failure is visible. Retry, notification, and history behavior follow the final failure model.	High	Blocked (Gap 3: failure and status)
35	Report and Workflow Dependency Connect a report to one or more Workflows. Edit the report, run the Workflows, then delete the report and run them again.	No Workflow sends the wrong, old, or invalid report, and one Workflow does not affect another. Exact edit and delete behavior follows the final dependency rules.	High	Blocked (Gap 3: Workflow dependency)
36	Several Scheduled Runs at the Same Time Schedule several reports for the same time and trigger them together.	Each run is handled separately. Reports are not mixed, silently dropped, or overwritten.	Medium	Ready
37	Generated File Storage and Retention Generate a report, try to download it again later, and test what happens after the defined retention period.	Download behavior follows one clear storage and retention rule. Old files do not stay available longer than intended, and it is clear whether the user gets an old file or a new generation.	Medium	Blocked (Gap 3: file retention)
7. Data Volume and Output Quality				
38	CSV at the 500K Limit Generate an Operational Data Report at or close to 500K matching records.	Supported records are not silently lost, and a partial export is never shown as complete. A timing threshold can be added once Engineering defines one.	High	Partial (Performance target)
39	CSV Above the 500K Limit Generate more than 500K matching records manually and in a scheduled run that grows above the limit.	The system must not silently cut the file and present it as complete. The final policy must be consistent. If truncation is allowed, the warning and record selection rule must be clear.	Critical	Blocked (Gap 4: over limit behavior)
40	CSV Data Period Generate CSV reports from data that covers a long period and compare the rows with the selected source and scope.	The CSV represents one clearly defined data period. Users should not receive an unexpectedly unlimited data set.	High	Blocked (Gap 4: CSV data period)

#	Scenario	Expected Result	Priority	Status
41	Valid Report with No Results Create a valid report whose scope or filters return zero results, including a scheduled run.	Zero results are clearly different from a technical failure or missing data. The exact empty report and delivery behavior must be decided before final pass or fail.	Medium	Blocked (Zero result behavior)
42	CSV Format and Spreadsheet Safety Export values with commas, quotes, line breaks, non Latin text, and values starting with =, +, -, or @. Open the file in a common spreadsheet tool.	The CSV opens correctly, values are kept accurately, and user controlled text does not trigger spreadsheet formulas. The exact escaping rule should be confirmed with Security.	High	Partial (Security: CSV escaping)
43	PDF Rendering Quality Generate a PDF from a complex Dashboard with many or wide widgets, long labels, page breaks, and a component that fails to render.	The PDF stays readable. There is no clipping, overlap, missing page, or silent missing content. A failed component is shown clearly.	High	Ready
8. Audit Logging				
44	Audit Log Coverage Create, update, delete, generate, download, and schedule report related actions, then inspect the audit log.	Create, update, and delete are logged as required by the PRD. Other actions follow the final audit scope, and the log must not suggest coverage it does not actually have.	Medium	Partial (Audit scope)

Release Acceptance Criteria

- All Critical tests pass, and no Critical test is still Blocked at release.
- No user, recipient, or tenant can access report data outside their allowed scope.
- PDF and CSV reports match the final source, scope, filters, fields, and time period.
- An incomplete, stale, failed, or unauthorized report is never shown or delivered as a valid complete report.
- Manual and scheduled generation and delivery work reliably, and failures are visible.
- Any open Critical or High decision from Part A is turned into a clear acceptance criterion and tested before release.

Explicitly Out of Scope

- Full testing of existing Workflow, RBAC, Dashboard, and Saved View functionality. Testing will focus only on their integration with Custom Reporting.
- Validation of the underlying security findings. The purpose of the test plan is to confirm that generated reports accurately reflect the data presented in the platform, rather than revalidate the original security data.
- Exhaustive testing of every possible field, filter, and module combination. Representative scenarios will be used, with additional module-specific coverage where behavior differs.
- Delivery methods other than email, as they are not included in the current release scope.
- Large-scale load and stress testing beyond the defined 500K-record scenario. This requires dedicated performance targets and separate performance testing.

Part C – BI & Monitoring Definition

Approach

Once Custom Reporting is live, the dashboard should answer two questions: are customers adopting and returning to the feature, and is it reliable enough for them to trust it? Adoption should be measured only against users who can create reports and customer accounts where the feature is enabled.

Because this is a new feature and the PRD does not define a production baseline, the alert thresholds below should be treated as starting points and adjusted after the first few weeks of real usage.

1. Product and Usage Monitoring

Customer Adoption Rate

Measures how broadly the feature is being adopted across customers, rather than being used heavily by only a few accounts.

Calculation: Customer accounts that generated at least one report during the period / customer accounts where Custom Reporting is enabled.

Report Creation Completion Rate

Shows whether users who start creating a report are able to finish the flow. A drop can point to unclear configuration, validation, or scope selection.

Calculation: Report creation flows that end with a saved report / report creation flows started.

Repeat and Scheduled Usage

A report that is reused or connected to a Workflow is a stronger sign of ongoing value than a one-time export. I would track both repeat generation and scheduled use.

Calculation: Saved reports generated more than once / saved reports generated at least once, and saved reports connected to an active Workflow / total saved reports.

Usage Mix

Shows how customers are using the feature in practice: Summary Report versus Operational Data, manual versus scheduled generation, and standard creation versus Save as Custom Report. These are supporting metrics rather than primary KPIs, but they can guide future product investment.

Calculation: Generations or creation flows grouped by report type, generation method, and entry point as a share of the total.

2. Quality and Reliability Monitoring

Report Generation Success Rate

The basic health metric for the feature. PDF and CSV should be viewed separately because the two report types can fail for different reasons.

Calculation: Successful report generations / total generation attempts.

Scheduled Delivery Success Rate

A scheduled report is only successful if it is both generated and delivered. This is especially important because no user may be present to notice a failure.

Calculation: Scheduled runs delivered successfully / scheduled runs that were due.

Report Generation Time

Shows whether the feature is becoming slow for normal use or for larger reports. I would monitor both typical performance and unusually slow cases.

Calculation: Time from generation start until the file is ready, tracked using the median and 95th percentile and split by report type.

Data Completeness and Source Failures

Two risks from Part A are worth monitoring directly: CSV reports reaching the 500K limit, and saved reports failing because their Dashboard, Saved View, or referenced field is no longer available.

Calculation: 500K-limit cases / CSV generations, and source-related failures / report generation attempts.

Generation and delivery failures should also be grouped by a small set of clear reasons such as permissions, missing source, size limit, timeout, or internal error. The overall failure rate shows that there is a problem; the reason helps identify where to investigate.

3. Alerting

Immediate alerts

The team should be notified automatically when there is a clear operational or security issue, for example:

- Report generation success falls below 95% for a sustained period, with enough attempts for the number to be meaningful.
- Scheduled delivery success falls below 95% over an hour, or repeated failures affect the same customer.
- There is any confirmed cross-tenant access or report data is exposed outside the authorized scope.
- Generation time increases significantly compared with the normal baseline for a sustained period.

The exact thresholds should be tuned after launch once normal traffic and performance patterns are known.

Weekly review

Adoption, creation completion, repeat and scheduled usage, report-type mix, 500K-limit cases, and the most common failure reasons are better reviewed weekly. These metrics mainly support product decisions and trend analysis, so they do not require an immediate alert.

4. Data We Would Need to Start Collecting

The PRD currently mentions audit logging for report creation, updates, and deletion, but that is not enough to calculate the metrics above. The product would need to record the main steps in the report journey: when creation starts and completes, where it started from, whether a report is generated manually or on a schedule, and whether it is later connected to a Workflow.

For generation and delivery, each run should record its result, duration, report type, row count or file size when relevant, and a clear failure reason. It should also be possible to identify cases where the 500K limit is reached, a source is missing, a permission check is rejected, or a scheduled delivery fails.

Two pieces of reference data are also needed: which customer accounts have the feature enabled and which users are allowed to create reports. Without this, adoption would be measured against the wrong population.

The analytics data should contain IDs and operational metadata only. It should not store report contents, finding details, report names, filter values, resource names, or recipient email addresses.